

# Gabriele Sarti

Postdoctoral Researcher, Northeastern University

Khoury College of Computer Sciences

02115 Boston, MA, USA

🏠 <https://gsarti.com>

✉ [gabriele.sarti996@gmail.com](mailto:gabriele.sarti996@gmail.com) | [g.sarti@northeastern.edu](mailto:g.sarti@northeastern.edu)

## CURRENT POSITION

---

### **Bau Lab, Northeastern University**

Boston, MA, USA

*Postdoctoral Researcher*

*Jan. 2026 – exp. Dec. 2026*

Advisor: David Bau. R&D for the NSF [National Deep Inference Fabric \(NDIF\)](#) project.

## EDUCATION

---

### **University of Groningen**

Groningen, Netherlands

*Ph.D. in Natural Language Processing, Cum Laude*

*Sept. 2021 – Dec. 2025*

Advisors: Arianna Bisazza, Malvina Nissim, Grzegorz Chrupala

Thesis: “[From Insights to Impact: Actionable Interpretability for Neural Machine Translation](#)”

### **University of Trieste & SISSA**

Trieste, Italy

*M.Sc. in Data Science and Scientific Computing, 110 Cum Laude*

*Oct. 2018 – Dec. 2020*

Advisors: Felice Dell'Orletta, Davide Crepaldi

Thesis: “[Interpreting Neural Language Models for Linguistic Complexity Assessment](#)”

### **Cégep de Saint-Hyacinthe**

Saint-Hyacinthe, QC, Canada

*Collegial Studies Degree (DEC) in Management Informatics*

*Aug. 2015 – May 2018*

Valedictorian of the 2018 informatics cohort. Program taught in French.

## EXPERIENCE

---

### INDUSTRIAL EXPERIENCE

#### **Amazon Web Services AI Lab**

New York, NY, USA

*Applied Scientist Intern, Translate (hosted by Georgiana Dinu, Maria Nadejde)*

*Jun. 2022 – Sept. 2022*

RAG-augmented LLM prompting for attribute-controlled translation.

#### **Aindo**

Trieste, Italy

*Research Scientist, Generative AI Systems*

*Nov. 2020 – Aug. 2021*

Structured clinical QA, recommender systems and sketch-to-image generation.

#### **Skytech Communications**

Montréal, QC, Canada

*Machine Learning Engineer Intern*

*Feb. 2018 – Jun. 2018*

Handwritten OCR detection and sentiment analysis for educational applications.

### ACADEMIC VISITING

#### **IRT Saint-Exupéry**

Toulouse, France

*Hosted by Fanny Jourdan*

*Feb. 2025 – Mar. 2025*

Collaboration on an open-source toolkit for concept-based interpretability of LLMs.

#### **Institute of Computational Linguistics (ILC-CNR)**

Pisa, Italy

*Hosted by Felice Dell'Orletta*

*Sept. 2019 – Dec. 2019*

Training NLP models for eye-tracking metrics prediction.

## TEACHING AND ADVISING

---

### CLASSES

Advanced Topics in Natural Language Processing, RUG, MSc Inf. Sci. (Co-instructor), 2025  
Natural Language Processing, RUG, MSc Inf. Sci. (Co-instructor), 2022-2024  
Fundamentals of Machine Learning, RUG, Data Science Minor (Instructor), 2024

### PH.D. STUDENTS

Sara Candussio, ADSAI Ph.D. at University of Trieste (co-sup. with Luca Bortolussi), 2024–  
Timothy Obiso, CBAI '26 Summer Fellow (co-sup. with Logan Riggs), 2026–

### MSC. STUDENTS

Jacopo Zacchigna, SPAR Spring '26 Fellow, 2026  
Hengxu Li, SPAR Spring '26 Fellow, 2026  
Khondoker I. Islam, NLD Project, University of Groningen, 2025  
Samuele D'Avenia, Thesis in Data Science, University of Trieste, 2024  
Sara Candussio, Thesis in Data Science, University of Trieste, 2024  
Konstantin Chernyshev, Thesis in NLP, University of Saarland & University of Groningen, 2024  
Daniel Scalena, Project in Computer Science, University of Milano-Bicocca, 2023  
Ludwig Sickert, Thesis in Artificial Intelligence, University of Groningen, 2023  
Qiankun Zheng, NLP Project, University of Groningen, 2023

### BSC. STUDENTS

Rohan Kathuria, CBAI '26 Summer Fellow (co-sup. with Logan Riggs), 2026–  
Aftab Iyigun, SPAR Spring '26 Research Project, 2026  
Mika Okamoto, SPAR Spring '26 Research Project, 2026  
Vinit Patel, SPAR Spring '26 Research Project, 2026

### AWARDS

---

**Best PhD Thesis Award**, European Association for Machine Translation (EAMT), 2026  
**Best Paper**, Italian Conference on Computational Linguistics (CLiC-it 2025), 2025  
**Best Paper**, Italian Conference on Computational Linguistics (CLiC-it 2023), 2023  
**Best Master's Thesis Award**, Italian Association for Computational Linguistics (AILC), 2022  
**Outstanding Reviewer**, EMNLP 2024, 2024

### SCHOLARSHIPS AND GRANTS

---

**Research Grant for Tinker API Credits** (USD 5k), Thinking Machines, 2026  
**Research Grant for Pangram API Credits** (USD 5k), Pangram, 2026  
**Research Grant for Compute Credits** (USD 5k), Google Cloud, 2026  
**Grant for Project Mentorship** (USD 7k), Supervised Program for Alignment Research, 2026  
**AI-net Fellowship** (Visiting expenses), Deutscher Akademischer Austauschdienst (DAAD), 2025  
**Language Technology Grant** (EUR 20k), Imminent Research Center, 2023  
**eScience Fellowship** (EUR 2k), Netherlands eScience Center, 2023  
**InDeep Grant** (PhD salary & research budget), Dutch Research Council (NWO), 2021  
**Data Science Excellence Fellowship** (EUR 12k), SISSA, 2018  
**Informatics Merit Scholarship** (EUR 1k), Hydro-Québec, 2016

**Journal Articles**

- [J6] Malvina Nissim<sup>†</sup>, Danilo Croce<sup>†</sup>, Viviana Patti<sup>†</sup>, Pierpaolo Basile<sup>†</sup> et al. (incl. Gabriele Sarti). “[Challenging the Abilities of Large Language Models in Italian: a Community Initiative](#)”. In *Italian J. of Computational Linguistics (IJCoL)*. Issue 12-1, 2026. (<sup>†</sup> = Equal contribution)
- [J5] Center for AI Safety, Scale AI, Humanity's Last Exam Contributors (incl. Gabriele Sarti). “[A benchmark of expert-level academic questions to assess AI capabilities](#)”. In *Nature*. 649, 1139–1146, 2026.
- [J4] Gabriele Sarti, Vilém Zouhar, Grzegorz Chrupała, Ana Guerberofo-Arenas, Malvina Nissim, Arianna Bisazza. “[QE4PE: Word-level Quality Estimation for Human Post-Editing](#)”. In *Transactions of the ACL (TACL)*. Issue 13: 1410–1435, 2025.
- [J3] Lukas Edman, Gabriele Sarti, Antonio Toral, Gertjan van Noord, Arianna Bisazza. “[Are Character-level Translations Worth the Wait? Comparing ByT5 and mT5 for Machine Translation](#)”. In *Transactions of the ACL (TACL)*. Issue 12: 392–410, 2024.
- [J2] Alessio Miaschi, Gabriele Sarti, Dominique Brunato, Felice Dell'Orletta, Giulia Venturi. “[Probing Linguistic Knowledge in Italian Neural Language Models across Language Varieties](#)”. In *Italian J. of Computational Linguistics (IJCoL)*. Issue 8-1, 2022.
- [J1] Ginevra Carbone, Gabriele Sarti. “[ETC-NLG: End-to-end Topic-Conditioned Natural Language Generation](#)”. In *Italian J. of Computational Linguistics (IJCoL)*. Issue 6-2, 2020.

**Conference Proceedings**

- [C13] Raghu Arghal<sup>†</sup>, Fade Chen<sup>†</sup>, Niall Dalton<sup>†</sup>, Evgenii Kortukov<sup>†</sup>, Calum McNamara<sup>†</sup>, Angelos Nalmpantis<sup>†</sup>, Moksh Nirvaan<sup>†</sup>, Gabriele Sarti, Mario Giulianelli. “[A Behavioural and Representational Evaluation of Goal-Directedness in Language Model Agents](#)”. In *International Conference on Machine Learning (ICML)*. 2026. (<sup>†</sup> = Alphabetical order)
- [C12] Antonin Poché<sup>†</sup>, Thomas Mullor, Gabriele Sarti et al., Fanny Jourdan<sup>†</sup>. “[Interpreto: An Explainability Library for Transformers](#)”. In *Annual Meeting of the ACL: System Demonstrations (ACL Demo)*. 2026. (<sup>†</sup> = Equal contribution)
- [C11] Daniel Scalena<sup>†</sup>, Gabriele Sarti<sup>†</sup>, Arianna Bisazza, Elisabetta Fersini, Malvina Nissim. “[Steering Large Language Models for Machine Translation Personalization](#)”. In *Conference of the European Chapter of the ACL (EACL)*. 2026. (<sup>†</sup> = Equal contribution)
- [C10] Gabriele Sarti, Vilém Zouhar, Malvina Nissim, Arianna Bisazza. “[Unsupervised Word-level Quality Estimation for Machine Translation Through the Lens of Annotators \(Dis\)agreement](#)”. In *Empirical Methods in Natural Language Processing (EMNLP)*. 2025. **Top 15%**.
- [C9] Sara Candussio, Gaia Saveri, Gabriele Sarti, Luca Bortolussi. “[Bridging Logic and Learning: Decoding Temporal Logic Embeddings via Transformers](#)”. In *European Conference on Machine Learning and Knowledge Discovery in Databases (ECML-PKDD)*. 2025.
- [C8] Mohammad Reza Ghasemi Madani, Aryo Pradipta Gema, Yu Zhao, Gabriele Sarti, Pasquale Minervini, Andrea Passerini. “[Noiser: Bounded Input Perturbations for Attributing Large Language Models](#)”. In *Conference on Language Modeling (COLM)*. 2025.
- [C7] Jirui Qi<sup>†</sup>, Gabriele Sarti<sup>†</sup>, Raquel Fernández, Arianna Bisazza. “[Model Internals-based Answer Attribution for Trustworthy Retrieval-Augmented Generation](#)”. In *Empirical Methods in Natural Language Processing (EMNLP)*. 2024. (<sup>†</sup> = Equal contribution)

- [C6] Gabriele Sarti, Grzegorz Chrupała, Malvina Nissim, Arianna Bisazza. “[Quantifying the Plausibility of Context Reliance in Neural Machine Translation](#)”. In *International Conference on Learning Representations (ICLR)*. 2024.
- [C5] Anna Langedijk, Hosein Mohebbi, Gabriele Sarti, Willem Zuidema, Jaap Jumelet. “[DecoderLens: Layerwise Interpretation of Encoder-Decoder Transformers](#)”. In *Findings of the Association for Computational Linguistics (NAACL)*. 2024.
- [C4] Gabriele Sarti, Malvina Nissim. “[IT5: Text-to-text Pretraining for Italian Language Understanding and Generation](#)”. In *Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)*. 2024.
- [C3] Gabriele Sarti, Nils Feldhus, Ludwig Sickert, Oskar van der Wal, Malvina Nissim, Arianna Bisazza. “[Inseq: An Interpretability Toolkit for Sequence Generation Models](#)”. In *Annual Meeting of the ACL: System Demonstrations (ACL Demo)*. 2023.
- [C2] Gabriele Sarti, Phu Mon Htut, Xing Niu, Benjamin Hsu, Anna Currey, Georgiana Dinu, Maria Nadejde. “[RAMP: Retrieval and Attribute-Marking Enhanced Prompting for Attribute-Controlled Translation](#)”. In *Annual Meeting of the ACL*. 2023.
- [C1] Gabriele Sarti, Arianna Bisazza, Ana Guerberof-Arenas, Antonio Toral. “[DivEMT: Neural Machine Translation Post-Editing Effort Across Typologically Diverse Languages](#)”. In *Empirical Methods in Natural Language Processing (EMNLP)*. 2022.

## Workshop Proceedings

- [W17] Cristiano Ciaccio, Asya Zanollo, Alessio Miaschi, Gabriele Sarti, Felice Dell'Orletta, Malvina Nissim. “[Reading Between the Clues: A Fine-Grained Analysis of Italian Crossword Solving Systems](#)”. In *Proceedings of the 12th Italian Conference on Computational Linguistics (CLiC-it)*. 2026.
- [W16] Sara Candussio<sup>†</sup>, Daniel Scalena<sup>†</sup>, Elisabetta Fersini, Luca Bortolussi, Malvina Nissim, Gabriele Sarti. “[Demystifying Entropy-based Selection for Chain-of-Thought Compression in Large Reasoning Models](#)”. In *Proceedings of the 9th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*. 2026. (<sup>†</sup> = Equal contribution)
- [W15] Sara Candussio, Gabriele Sarti, Gaia Saveri, Luca Bortolussi. “[Distilling Formal Logic into Neural Spaces: A Kernel Alignment Approach for Signal Temporal Logic](#)”. In *International Conference on Neurosymbolic Learning and Reasoning (NeSy)*. 2026.
- [W14] Stefano Atzeni, Gabriele Sarti, Tommaso Caselli, Malvina Nissim. “[TiGRO at FadeIT: E Pluribus Unum - A Multi-task Approach to Fallacy Detection and Span Identification](#)”. In *9th Evaluation Campaign of Natural Language Processing Tools for Italian (EVALITA)*. 2026.
- [W13] Cristiano Ciaccio, Gabriele Sarti, Alessio Miaschi, Felice Dell'Orletta, Malvina Nissim. “[Cruciverb-IT@ EVALITA 2026: Overview of the Crossword Solving in Italian Task](#)”. In *9th Evaluation Campaign of Natural Language Processing Tools for Italian (EVALITA)*. 2026.
- [W12] Khondoker Ittehadul Islam, Gabriele Sarti. “[Reveal-Bangla: A Dataset for Cross-Lingual Multi-Step Reasoning Evaluation](#)”. In *Second Workshop on Bangla Language Processing (BLP)*. 2025.
- [W11] Dana Arad, Yonatan Belinkov, Hanjie Chen, Najoung Kim, Hosein Mohebbi, Aaron Mueller, Gabriele Sarti, Martin Tutek. “[Findings of the BlackboxNLP 2025 Shared Task: Localizing Circuits and Causal Variables in Language Models](#)”. In *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*. 2025.

- [W10] Cristiano Ciaccio, Gabriele Sarti, Alessio Miaschi, Felice Dell'Orletta. “[Crossword Space: Latent Manifold Learning for Italian Crosswords and Beyond](#)”. In *Proceedings of the 11th Italian Conference on Computational Linguistics (CLiC-it)*. 2025. **Best paper award**.
- [W9] Gabriele Sarti, Tommaso Caselli, Malvina Nissim, Arianna Bisazza. “[Non Verbis, Sed Rebus: Large Language Models Are Weak Solvers of Italian Rebuses](#)”. In *Proceedings of the 10th Italian Conference on Computational Linguistics (CLiC-it)*. 2024.
- [W8] Gabriele Sarti, Tommaso Caselli, Arianna Bisazza, Malvina Nissim. “[EurekaRebus - Verbalized Rebus Solving with LLMs: A CALAMITA Challenge](#)”. In *Proceedings of the Tenth Italian Conference on Computational Linguistics (CLiC-it 2024)*. 2024.
- [W7] Daniel Scalena, Gabriele Sarti, Malvina Nissim. “[Multi-property Steering of Large Language Models with Dynamic Activation Composition](#)”. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*. 2024.
- [W6] Gabriele Sarti, Nils Feldhus, Jirui Qi, Malvina Nissim, Arianna Bisazza. “[Democratizing Advanced Attribution Analyses of Generative Language Models with the Inseq Toolkit](#)”. In *2nd World Conference on Explainable AI (xAI): System Demonstrations*. 2024.
- [W5] Federico Bianchi, Giuseppe Attanasio, Raphael Pisoni, Silvia Terragni, Gabriele Sarti, Dario Balestri. “[Contrastive Language-Image Pre-training for the Italian Language](#)”. In *Proceedings of the 9th Italian Conference on Computational Linguistics (CLiC-it)*. 2023. **Best paper award**.
- [W4] Gabriele Sarti, Dominique Brunato, Felice Dell'Orletta. “[That Looks Hard: Characterizing Linguistic Complexity in Humans and Language Models](#)”. In *Workshop on Cognitive Modeling and Computational Linguistics (CMCL)*. 2021. **Best paper award runner-up**.
- [W3] Ludovica Pannitto, Lucia Busso, Claudia Roberta Combei, Lucio Messina, Alessio Miaschi, Gabriele Sarti, Malvina Nissim. “[Teaching NLP with Bracelets and Restaurant Menus: An Interactive Workshop for Italian Students](#)”. In *Fifth Workshop on Teaching NLP*. 2021.
- [W2] Alessio Miaschi, Gabriele Sarti, Dominique Brunato, Felice Dell'Orletta, Giulia Venturi. “[Italian Transformers Under the Linguistic Lens](#)”. In *Proceedings of the 7th Italian Conference on Computational Linguistics (CLiC-it)*. 2020. **Best paper award runner-up**.
- [W1] Gabriele Sarti. “[Improving Complexity and Acceptability Prediction with Multi-task Learning on Self-Supervised Annotations](#)”. In *Proceedings of the Seventh Evaluation Campaign of Natural Language Processing and Speech Tools for Italian (EVALITA)*. 2020. **Best system award**.

## Preprints

- [P6] Vilém Zouhar, Niyati Bafna, Mukund Choudhary, Maike Züfle, Sara Rajae, Pinzhen Chen et al. (incl. Gabriele Sarti). “[Last Translation Benchmark](#)”. In *CoRR/2609.04173*. 2026.
- [P5] Daniel Scalena<sup>†</sup>, Sara Candussio<sup>†</sup>, Luca Bortolussi, Elisabetta Fersini, Malvina Nissim, Gabriele Sarti. “[Beyond the Commitment Boundary: Probing Epiphenomenal Chain-of-Thought in Large Reasoning Models](#)”. In *CoRR/2606.13603*. 2026. (<sup>†</sup> = **Equal contribution**)
- [P4] Evgenii Kortukov, Piotr Komorowski, Florian Klein, Paula Engl, Gabriele Sarti, Seong Joon Oh, Sebastian Lapuschkin, Wojciech Samek. “[Predicting Future Behaviors in Reasoning Models Enables Better Steering](#)”. In *CoRR/2606.11172*. 2026.
- [P3] Vera Neplenbroek, Gabriele Sarti, Arianna Bisazza, Raquel Fernández. “[Topics as Proxies for Sociodemographics: How Conversational Context Affects LLM Answers](#)”. In *CoRR/2606.02776*. 2026.
- [P2] Natalie Shapira, Chris Wendler, Avery Yen, Gabriele Sarti et al., David Bau. “[Agents of Chaos](#)”. In *CoRR/2602.20021*. 2026.

- [P1] Javier Ferrando, Gabriele Sarti, Arianna Bisazza, Marta R. Costa-jussà. “[A Primer on the Inner Workings of Transformer-based Language Models](#)”. In *CoRR/2405.00208*. 2024.

## Manuscripts

- [M2] Gabriele Sarti. “[From Insights to Impact: Actionable Interpretability for Neural Machine Translation](#)”. *PhD Thesis, Univ. of Groningen*. 2025. **Cum Laude, EAMT Best Thesis Award**.
- [M1] Gabriele Sarti. “[Interpreting Neural Language Models for Linguistic Complexity Assessment](#)”. *MSc. Thesis in Data Science, University of Trieste*. 2020. **AILC Best Master's Thesis Award**.

## Edited Works

- [E1] Yonatan Belinkov, Aaron Mueller, Najoung Kim, Hosein Mohebbi, Hanjie Chen, Dana Arad, Gabriele Sarti. “[Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP](#)”. 2025.

## PROFESSIONAL SERVICE

---

### Organizing Committee

EMNLP BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP, 2025 – 2026  
NDIF [Aletheia's Quest](#) Competition on LLM Lie Detection, 2026  
EVALITA Shared Task on Automatic Italian Crossword Solving (Cruciverb-IT), 2026

### Area Chair

AAAI Conference on Artificial Intelligence (AAAI), Alignment Track, 2027  
ICML/NeurIPS Workshop on Mechanistic Interpretability, 2026-  
Conference on Language Modeling (COLM), 2026  
Annual Meeting of the Association for Computational Linguistics (ACL, via ARR), 2025 – 2026  
Empirical Methods for Natural Language Processing (EMNLP, via ARR), 2025  
Italian Conference on Computational Linguistics (CLiC-it), 2025

### Academic Paper Reviewer

COLM Actionable Interpretability Workshop (AIW), 2026  
Neural Information Processing Systems (NeurIPS), 2026-  
Transactions on Machine Learning Research (TMLR), 2025-  
International Conference on Learning Representations (ICLR), 2025 – 2026  
EMNLP BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP, 2023 – 2025  
Annual Conference of the European Chapter of ACL (EACL, via ARR), 2025  
Annual Conference of the Asia-Pacific Chapter of ACL (AACL, via ARR), 2025  
ICML/NeurIPS Workshop on Mechanistic Interpretability, 2024  
Conference on Language Modeling (COLM), 2024  
Annual Conference of the Nations of the Americas Chapter of ACL (NAACL, via ARR), 2024  
Empirical Methods for Natural Language Processing (EMNLP, via ARR), 2022 – 2024  
Italian Conference on Computational Linguistics (CLiC-it), 2021, 2024, 2025  
Annual Meeting of the Association for Computational Linguistics (ACL, via ARR), 2023  
International Conference on Language Resources and Evaluation (LREC), 2022

### Grant Reviewer

Schmidt Sciences, Interpretability RFP (Stage 1), 2026

### INVITED LECTURES

#### **Interpretability for Language Models: Current Trends and Applications**

University of Pisa, Mechanistic Interpretability of LLMs (Ph.D., host: Andrea Esuli), 2026.  
University College London, Adv. Computational Linguistics (MSc., host: Mario Giulianelli), 2026.  
University of Trieste, Explainable and Neurosymbolic AI (MSc., host: Luca Bortolussi), 2025.  
University of Groningen, Trustworthy and Explainable AI (MSc., host: Marco Zullo), 2025.  
Sapienza University of Rome, Explainable AI (Ph.D., host: Simone Scardapane), 2024.

#### **Interpreting Context Usage in Language Models**

Northeastern University, [CS7810: Neural Mechanics](#) (MSc., host: David Bau), 2026.  
Politecnico di Torino, Explainable and Trustworthy AI (MSc., host: Eliana Pastor), 2024.  
University of Pisa, AILC Lectures on Computational Linguistics (with Alessio Miaschi), 2023.

### ACADEMIC SEMINARS

#### **Scaling Interpretability for LLM Agents**

BauLab, Northeastern University. Boston, MA, USA. 2026.  
Algorithmic Alignment Group, MIT CSAIL. Boston, MA, USA. 2026

#### **Interpreting Context Usage in Generative Language Models**

MT Group, Fondazione Bruno Kessler (FBK). Online. 2026.  
Mueller Lab, Boston University. Boston, MA, USA. 2026.  
Language Technology Group, University of Helsinki. Online. 2026.  
Dept. of Engineering, University of Padua. Padua, IT. 2025.  
LanD Group, Fondazione Bruno Kessler (FBK). Online. 2025.  
ANITI and DEEL Groups, IRT Saint-Exupéry. Toulouse, FR. 2025.  
Center for Information and Language Processing (CIS), LMU Munich. Munich, DE. 2024.  
Laboratory of Data Engineering (LADE), Area Science Park. Trieste, IT. 2024.  
SheffieldNLP, University of Sheffield. Online. 2024.  
Netherlands eScience Center. Amsterdam, NL. 2023.  
REST-CL, Pompeu Fabra University. Tarragona, ES. 2023.  
AI-Lab, University of Trieste. Trieste, IT. 2023.  
GroNLP, University of Groningen. Groningen, NL. 2023.

### OTHER PRESENTATIONS

University of Groningen, Interpreting Latent Features in Large Language Models. Groningen. 2025.  
DFKI Saarbrücken, QE4PE: Word-level Quality Estimation for Human Post-Editing. Online. 2025.  
Sapienza University of Rome, An Interpretability Toolkit for Language Models. Rome. 2023.  
Radboud University, Advanced Interpretability for Language Models. Nijmegen. 2023.  
NeurIPS XAI4Debugging WS, Interpretable Interactive Neural Machine Translation. Online. 2021.  
Bocconi University, Linguistic Complexity in Humans and Language Models. Online. 2021.  
University of Trieste, Language Models: the New Frontier of Language Understanding. Trieste. 2019.  
38th AQPC Symposium, The Educational Impact of Artificial Intelligence. Saint-Hyacinthe. 2018.

### SCIENCE COMMUNICATION

The Inquisitive Mind Magazine, Peer reviewer. 2025.  
Trieste Next, Inside the Algorithm: Transparency and Impact of Generative AI. Trieste. 2025.  
[AI2S](#), Aprire la scatola nera dei modelli del linguaggio: rischi e opportunità. Trieste. 2024.  
Trieste Next, AI-Italo Svevo: Lettere da un'intelligenza artificiale. Live demo. Trieste. 2019.  
[Trieste Science Fiction Festival](#), The Literary Ordnance: When the Writer is an AI. Trieste. 2019.

### INDUSTRIAL SEMINARS

InDeep Masterclass, Interpreting LLMs and Other Deep Learning Models. Deloitte Amsterdam. 2025.  
InDeep Masterclass, Explaining Foundation Models. University of Amsterdam. 2023.

Translated S.r.l. Seminar, Towards User-centric Interpretability of NLP Models. Online. 2022.

## PRESS COVERAGE

---

**The Sidechannel**, ‘Abliterated’ AI model lower the bar for cyberattacks, experts warn. 2026.  
**Wired**, OpenClaw Agents Can Be Guilt-Tripped Into Self-Sabotage. 2026.  
**Science**, AI algorithms can become 'agents of chaos'. 2026.  
**Die Zeit**, Das ist erst der Anfang des Chaos. 2026.  
**The Times of India**, When AI gets keys, 'agents of chaos' leak secrets, wipe systems. 2026.  
**TG3 Mondo**, Il Vaso di Pandora - Armi autonome, AI e Autocrazie. 2026.  
**Slator**, Italian Benchmark Evaluates Large Language Models, Includes AI Translation. 2026.  
**Slator**, Does Word-Level Quality Estimation Really Improve AI Translation Post-Editing? 2025.  
**MarkTechPost**, Deciphering Language Models: Advances in Interpretability Research. 2024.  
**Bocconi University Hub News**, L'AI che racconta le immagini in italiano. 2021.  
**Poliflash**, CLIP-Italian: un nuovo modello di AI per collegare immagini e testi in italiano. 2021.

## MENTORSHIP AND SOCIAL ENGAGEMENTS

---

|             |                                                                                               |
|-------------|-----------------------------------------------------------------------------------------------|
| Summer 2026 | Mentor, <a href="#">Cambridge Boston Alignment Initiative (CBAI) Fellowship</a>               |
| Spring 2026 | Mentor, Implicit Personalization, Supervised Program for Alignment (SPAR)                     |
| Fall 2025   | Co-mentor for Project Telos, <a href="#">Supervised Program for Alignment (SPAR)</a>          |
| 2022 - 2025 | Research and Upskilling Advisor, <a href="#">Safe AI Netherlands (SAIN)</a> Groningen         |
| Winter 2021 | Team Mentor, <a href="#">PiCampus School of AI</a>                                            |
| 2020        | Co-founder and President, <a href="#">Artificial Intelligence Student Society (AI2S)</a>      |
| 2020 - 2021 | SciComm Team Member, <a href="#">Italian Association for Computational Linguistics (AILC)</a> |

## ACADEMIC REFERENCES

---

**Mario Giulianelli**, Associate Professor, University College London — [m.giulianelli@ucl.ac.uk](mailto:m.giulianelli@ucl.ac.uk)  
**David Bau**, Assistant Professor, Northeastern University — [d.bau@northeastern.edu](mailto:d.bau@northeastern.edu)  
**Malvina Nissim**, Full Professor, University of Groningen — [m.nissim@rug.nl](mailto:m.nissim@rug.nl)  
**Arianna Bisazza**, Associate Professor, University of Groningen — [a.bisazza@rug.nl](mailto:a.bisazza@rug.nl)  
**Luca Bortolussi**, Full Professor, University of Trieste — [lbortolussi@units.it](mailto:lbortolussi@units.it)